ABEL YAGUBYAN

abelyagubyan@berkeley.edu | abelo9996.github.io | GitHub | LinkedIn | Google Scholar | ORCID

EDUCATION

University of Georgia, Athens: Ph.D. Artificial Intelligence **August 2026 - Present**
Research focus: trustworthy AI evaluation, including the reliability of LLM-as-a-judge scoring and the reproducibility of LLM agents.

Northwestern University: M.S. Computer Science, Summa Cum Laude **September 2022 - June 2023**
Concentration in Deep Learning & Distributed Systems. Coursework: Deep Learning, Deep Learning in NLP, Deep Reinforcement Learning, CUDA Programming, Low-level/Kernel Software Development, Advanced Networking, Embedded Systems

University of California, Berkeley: B.A. Computer Science, B.A. Applied Mathematics **August 2018 - May 2022**
Coursework: AI, ML, Computer Vision, Parallel Computers, Computer Architecture, Operating Systems, Database Systems, Computer Security, Algorithms, Stochastic Optimization, Advanced Linear Algebra

PUBLICATIONS

- **Yagubyan, A.** *The Coin Flip Judge? Reliability and Bias in LLM-as-a-Judge Evaluation.* GroundLM Workshop at EMNLP 2026, Archival Long Paper (to appear). arXiv:2606.13685
- **Yagubyan, A.** *How Consistent Are LLM Agents? Measuring Behavioral Reproducibility in Multi-Step Tool-Calling Pipelines.* Preprint, 2026. arXiv:2605.28840
- Zheng, W.; Stahl, B. E.; Filippenko, A. V.; et al., incl. **Yagubyan, A.** *The Lick Observatory Supernova Search follow-up program: photometry data release of 70 SESNe.* Monthly Notices of the Royal Astronomical Society 512(3), 3195–3214, 2022. DOI
- **Yagubyan, A.**; Garcia, D. *Seamless Embedding of Programming IDEs into Computer-Based Testing Software.* ACM SIGCSE Technical Symposium (Vol. 2), March 2022. DOI

RESEARCH EXPERIENCE

Lawrence Berkeley National Laboratory: Pagoda Project Group
Research Contributor *June 2022 – April 2023*
- Developed and deployed an automated performance regression testing platform for UPC++, a C++ library for writing efficient and scalable parallel programs on distributed-memory parallel computers, using C++, UPC++, Python, and Shell scripts.
- Implemented a distributed hashing benchmark in UPC++ and compared its performance against other parallel-programming libraries, such as MPI and SHMEM.

SCALPEL: Customized Deep Neural Network Compression
Research Project, Northwestern University *January 2022 – June 2022*
- Built a Python DNN compression framework that tailors pruning to the parallelism of the target hardware (CPU, GPU, microcontroller), **reducing model size by up to 80%** and **achieving a 3.5x speedup** on AlexNet and LeNet-5.

OPEN SOURCE

TruLens (Snowflake): LLM and agent evaluation library
Area Triager *September 2026 – Present*
- Listed in MAINTAINERS.md with scope over OpenTelemetry instrumentation, evaluation and database correctness, and shared core code.
- **13 merged pull requests**: streaming-token, time-to-first-token, and throughput metrics in the OpenTelemetry path, plus twelve correctness fixes across persistence, dashboard, evaluation API, IR metrics, OpenAI and Anthropic providers, and leaderboard aggregation.
- Proposed repeated-trial self-consistency for LLM-as-a-judge metrics (#2741), grounded in the Coin Flip Judge paper.

inspect_evals (UK AI Security Institute)
Contributor *September 2026*
- Merged PR #2317 migrating the shared metric helpers (mean, standard deviation, standard error) to the Inspect aggregation API; credited in release v0.20.0.

LearnOS: open-source agentic AI university
Creator & Maintainer *2026 – Present*
- Self-hosted system of specialized agents (curriculum, Socratic tutor, assessment, research, analytics, profiling) that builds personalized learning roadmaps and grades toward mastery-based certification. Node.js, Vite, SQLite, OpenRouter.

INDUSTRY EXPERIENCE

Aivaro

Co-founder & CEO *March 2026 – June 2026*

- Founded an AI workflow execution platform that let small and medium-sized businesses automate revenue-critical operations across Gmail, Stripe, Google Sheets, Calendly, and similar tools.
- Architected a deterministic LLM-to-workflow system that converts natural language into structured, executable DAGs with built-in reliability (retries, monitoring, approval gating).
- Deployed in a commercial SMB environment, automating booking, payment, and reconciliation workflows, **saving 5 hours/week** of manual effort and **lowering no-shows from 20–25% to 8%**.
- Selected into Y Combinator's Spring 2026 batch (<2% acceptance rate).

C3.ai

Senior Data Scientist *May 2024 – March 2026*

- Led data science initiatives in predictive maintenance and generative AI across large-scale customer deployments (Shell, ExxonMobil, Dow, Bloom Energy). **Directed pilots delivering over $10M in value** and **deploying over 1,000 models** per project.
- Developed a foundational healthcare transformer model on EHR data for a customer deployment, **trained on 1M+ U.S. patient records and 1B+ anonymized datapoints**.
- Created *SpeedyREL*, an internal tool accelerating collaboration between data scientists and subject matter experts in deploying semi-supervised reconstruction-based predictive maintenance models. **Adopted by 5+ projects**, **reducing deployment timelines by up to 50%**, and incorporated into the onboarding and training materials for all of C3.ai's Reliability teams.
- Promoted to Senior Data Scientist in under 18 months for consistent delivery of high-impact customer solutions and contributions to core product development.

FibonAI

Co-founder *June 2023 – February 2024*

- Designed and wrote the source code for a SaaS platform providing an LLM-powered workspace for in-house corporate legal teams, integrating intake & triage, contract management, matter management, and AI assistants.
- **Onboarded 8 General Counsels/Chief Legal Officers** at large in-house corporate legal departments, plus 2 more customers to the waitlist through organic acquisition.
- Achieved a **60% conversion rate** in product demonstrations and **100% weekly user growth**.
- Accepted into UC Berkeley's SkyDeck Pad-13 Incubator Program, competing against 5,000+ startups globally.

Apple x UC Berkeley

Software Engineering Intern *May 2021 – August 2021*

- Project lead for a collaboration between Apple and UC Berkeley's CS61C (Computer Architecture), building and managing computer architecture, algorithms, and ML projects for **4,000+ students**, using C, C++, Python, SIMD, OpenMP, and RISC-V assembly.

TECHNICAL SKILLS

- **Languages:** Python, C++, C, CUDA, Rust, Go, Java, JavaScript, SQL, MATLAB, R
- **ML & LLM Evaluation:** PyTorch, TensorFlow, Hugging Face Transformers, XGBoost, Inspect AI, TruLens, OpenTelemetry
- **HPC & Systems:** MPI, UPC++, SHMEM, OpenMP, RISC-V, x86
- **Infrastructure:** Docker, Kubernetes, AWS, Google Cloud, MongoDB, MySQL, Node.js, React, Django, Flask, Bash

SERVICE

- **Peer reviewer, 200+ reviews across 9 Elsevier journals** (2024–2026), including *Engineering Applications of Artificial Intelligence*, *Image and Vision Computing*, *Neural Networks*, and *Information Fusion*. Elsevier Recognised Reviewer certificates for each journal.