ABEL YAGUBYAN

abelyagubyan@berkeley.edu | abelo9996.github.io | GitHub | LinkedIn | Google Scholar | ORCID

EDUCATION

University of Georgia, Athens: Ph.D. Artificial Intelligence **August 2026 - Present**
Research focus: reliability of LLM-as-a-judge evaluation and reproducibility of LLM agents.

Northwestern University: M.S. Computer Science, Summa Cum Laude **September 2022 - June 2023**
Coursework: Deep Learning, Deep Learning in NLP, Deep Reinforcement Learning, CUDA Programming, Kernel Development

University of California, Berkeley: B.A. Computer Science, B.A. Applied Mathematics **August 2018 - May 2022**
Coursework: AI, Machine Learning, Computer Vision, Parallel Computers, Operating Systems, Databases, Optimization

PUBLICATIONS

- **Yagubyan, A.** *The Coin Flip Judge? Reliability and Bias in LLM-as-a-Judge Evaluation.* GroundLM @ EMNLP 2026 (Archival). arXiv:2606.13685
- **Yagubyan, A.** *How Consistent Are LLM Agents? Measuring Behavioral Reproducibility in Multi-Step Tool-Calling Pipelines.* Preprint, 2026. arXiv:2605.28840
- Zheng, W.; Stahl, B. E.; Filippenko, A. V.; et al., incl. **Yagubyan, A.** *The Lick Observatory Supernova Search follow-up program: photometry data release of 70 SESNe.* MNRAS 512(3), 2022.
- **Yagubyan, A.**; Garcia, D. *Seamless Embedding of Programming IDEs into Computer-Based Testing Software.* ACM SIGCSE 2022.

OPEN SOURCE

TruLens (Snowflake), Area Triager **September 2026 – Present**
- Listed in MAINTAINERS.md; **13 merged PRs**, including streaming and time-to-first-token metrics in the OpenTelemetry path and twelve correctness fixes across persistence, evaluation API, providers, and leaderboard aggregation.
- Proposed repeated-trial self-consistency for LLM-as-a-judge metrics (#2741), grounded in the Coin Flip Judge paper.

inspect_evals (UK AI Security Institute), Contributor **September 2026**
- Migrated the shared metric helpers to the Inspect aggregation API (PR #2317); credited in release v0.20.0.

INDUSTRY EXPERIENCE

Aivaro
Co-founder & CEO, Y Combinator Spring 2026 *March 2026 – June 2026*
- Founded an AI workflow platform automating revenue-critical SMB operations across Gmail, Stripe, Google Sheets, and Calendly.
- Architected a deterministic LLM-to-workflow system that converts natural language into executable DAGs with built-in reliability (retries, monitoring, approval gating).
- Deployed for booking, payment, and reconciliation, **saving 5 hours/week** and **cutting no-shows from 20–25% to 8%**.

C3.ai
Senior Data Scientist *May 2024 – March 2026*
- Led predictive maintenance and generative AI deployments for Shell, ExxonMobil, Dow, and Bloom Energy, **delivering over $10M in value** and **deploying 1,000+ models** per project.
- Developed a foundational healthcare transformer **trained on 1M+ U.S. patient records and 1B+ datapoints**.
- Created *SpeedyREL*, an internal tool for deploying reconstruction-based predictive maintenance models; **adopted by 5+ projects**, **cut deployment time by up to 50%**, and became part of onboarding for all Reliability teams.
- Promoted to Senior Data Scientist in under 18 months.

FibonAI
Co-founder *June 2023 – February 2024*
- Built an LLM-powered workspace for in-house legal teams (intake, contracts, matter management, AI assistants); **onboarded 8 General Counsels/CLOs** with a **60% demo conversion rate** and **100% weekly user growth**.
- Accepted into UC Berkeley's SkyDeck Pad-13 incubator, competing against 5,000+ startups globally.

Apple x UC Berkeley
Software Engineering Intern *May 2021 – August 2021*
- Project lead building CS61C architecture and ML projects for **4,000+ students** in C, C++, SIMD, OpenMP, and RISC-V.

SKILLS & SERVICE

- **Languages:** Python, C++, C, CUDA, Rust, Go, Java, JavaScript, SQL, MATLAB, R
- **Tools:** PyTorch, TensorFlow, Hugging Face, XGBoost, Inspect AI, TruLens, MPI, UPC++, Docker, Kubernetes, AWS, GCP
- **Peer review:** 200+ reviews across 9 Elsevier journals, including *Engineering Applications of AI* and *Neural Networks*